

Real-World Use Case of Amazon Redshift Data API for SQL Querying Without Persistent Connections

Problem Statement:

A data engineering team wants to allow multiple applications to run SQL queries on their Amazon Redshift cluster **without managing long-lived database connections or provisioning client drivers**. This capability is crucial to support event-driven applications, serverless workflows, and scheduled automation **without the overhead of connection pooling**.

Objective:

To simulate and experience a real-world solution using the Amazon Redshift **Data API**, enabling SQL query execution directly via AWS CLI or Lambda without requiring JDBC/ODBC connections.

Technologies Used:

- Amazon Redshift Serverless
 - AWS Secrets Manager
 - Amazon Redshift Data API
 - AWS CLI
-

Solution Steps Implemented:

- ◊ **Step 1: Provision Redshift Serverless**
 - Created a Redshift Serverless **workgroup**: my-redshift-new-workgroup
 - Created a **namespace** with default database: dev
 - Configured the admin username and password
- ◊ **Step 2: Initial Authentication Issues Identified**
 - Attempted login using the default admin user and **faced error**: > FATAL: password authentication failed for user "admin"
 - Root cause: Redshift Serverless auto-manages admin credentials, which mismatch with custom Secrets Manager entries.

◊ Step 3: Created Custom User for Authentication

- Logged into Query Editor v2 using the default session
- Created a custom user:

```
CREATE USER datauser PASSWORD 'DataUserPassword123!';
GRANT USAGE ON SCHEMA public TO datauser;
GRANT SELECT ON ALL TABLES IN SCHEMA public TO datauser;
```

◊ Step 4: Created Table and Sample Data

```
CREATE TABLE sales (
    id INT,
    product VARCHAR(50),
    quantity INT,
    sale_date DATE
);

INSERT INTO sales VALUES
(1, 'Keyboard', 10, '2025-06-15'),
(2, 'Mouse', 5, '2025-06-16'),
(3, 'Monitor', 2, '2025-06-17');
```

◊ Step 5: Stored Credentials in Secrets Manager

- Created a new secret named: redshift-data-api-datauser
- Credentials stored:
 - username: datauser
 - password: DataUserPassword123!
- Captured full ARN:

arn:aws:secretsmanager:ap-south-1:535002864770:secret:redshift-data-api-datauser-AqzP8D

◆ Step 6: Executed SQL Query via AWS CLI

```
```bash
aws redshift-data execute-statement \
--workgroup-name my-redshift-new-workgroup \
--database dev \
--sql "SELECT * FROM sales;" \
--secret-arn arn:aws:secretsmanager:ap-south-
1:535002864770:secret:redshift-data-api-datauser-AqzP8D \
--region ap-south-1
```

- Captured query-id from the CLI response.

#### ◊ Step 7: Retrieved Results via CLI

```
aws redshift-data get-statement-result --id <query-id> --region ap-south-1
```

- Got expected structured output in JSON format from the sales table.
- 

#### Final Outcome:

- Simulated a **real-world serverless SQL querying** use case.
  - Avoided JDBC/ODBC drivers using **Redshift Data API**.
  - Managed secure credentials using **AWS Secrets Manager**.
  - Prepared architecture for **event-based, scheduled, or Lambda-powered** Redshift operations.
- 

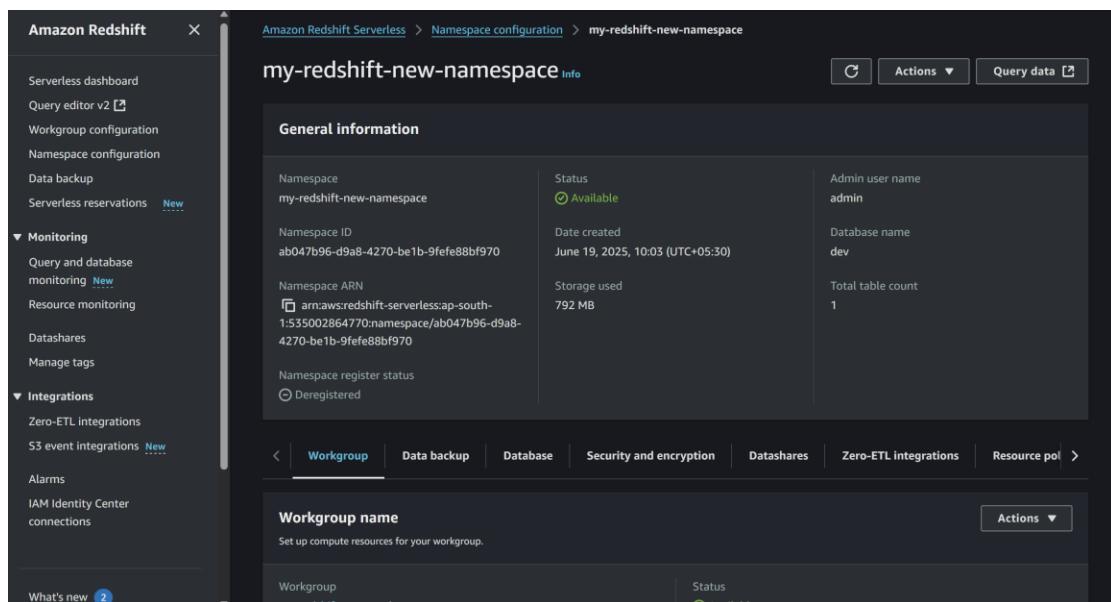
#### Next Steps (Optional Enhancements):

- Automate Data API + export results to CSV via Lambda
  - Attach IAM permission boundaries for more security
  - Build real-time dashboard with Streamlit or Flask using the Data API
- 

#### Workspace Output (Paste Your Results Below):

 Paste your screenshots, execution outputs, or CLI command results below:

[Redshift Workgroup Console Screenshot]



## [Secrets Manager ARN and structure]

The screenshot shows the AWS Secrets Manager interface for the 'redshift-data-api-datauser' secret. The 'Secret details' section includes:

- Encryption key: aws/secretsmanager
- Secret name: redshift-data-api-datauser
- Secret ARN: arn:aws:secretsmanager:ap-south-1:535002864770:secret:redshift-data-api-datauser-AqzP8D

The 'Secret value' section displays two entries:

Key/value	Secret value
username	D[REDACTED]@123!
password	D[REDACTED]123!

## [CLI output of execute-statement and get-statement-result]

```
C:\Users\7386847235>
C:\Users\7386847235>aws redshift-data get-statement-result --id fbe31709-78cb-4b2f-a842-00fela9783fb --region ap-south-1 --output table

| GetStatementResult
+-----+
| TotalNumRows | 3
+-----+
| ColumnMetadata
|-----+
| isCaseSensitive | isCurrency | isSigned | label | length | name | nullable | precision | scale | schemaName | tableName | type
eName
+-----+
| False | False | True | id | 0 | id | 1 | 10 | 0 | public | sales | in
t4
| True | False | False | product | 0 | product | 1 | 50 | 0 | public | sales | va
rchar
| False | False | True | quantityore | -- |
+-----+
```

---

## Summary

This project shows how to build **lightweight, scalable, and secure querying systems** using:

- Amazon Redshift Serverless
- AWS CLI
- Redshift Data API
- Secrets Manager

You've now implemented a **connectionless data querying model**, ready for real-world automation!

---

**Prepared by:** A. Siva Sandeep

**Date:** June 19, 2025

---

## Tags

#RedshiftDataAPI #AWSRedshiftServerless #SecretsManager #ServerlessSQL  
#CloudProject #EventDrivenSQL