

Optimizing Query Performance in Amazon Redshift by Resolving Data Skew

Project Objective:

To simulate and resolve uneven load distribution across nodes in an Amazon Redshift provisioned cluster by optimizing table distribution styles.

Problem Scenario:

A corporation using a Redshift provisioned cluster (4-node ra3.4xlarge) with DISTSTYLE EVEN observed **poor query performance during peak loads**. Monitoring revealed that one node consistently experienced **high CPU utilization**, while others were **underutilized**. This indicated **data skew** and poor slice-level balance.

The goal: **Redistribute data uniformly** across all nodes using appropriate distribution styles.

Phase 1: Problem Simulation

Cluster Setup

-  Created a Redshift provisioned cluster with **4 ra3.4xlarge nodes**

Dataset Selection

-  Used NYC Yellow Taxi trip data from **Kaggle**
-  Selected a small sample dataset (<50MB) to simulate skew behavior

S3 Upload

-  Uploaded CSV to:

s3://nyc-taxi-dataset-bucket/yellow_tripdata_2015-01.csv

Redshift Table Creation

-  Created table trips_even using DISTSTYLE EVEN:

```
CREATE TABLE trips_even (
    VendorID INT,
    tpep_pickup_datetime TIMESTAMP,
    tpep_dropoff_datetime TIMESTAMP,
```

```

    passenger_count INT,
    trip_distance FLOAT,
    pickup_longitude FLOAT,
    pickup_latitude FLOAT,
    RateCodeID INT,
    store_and_fwd_flag VARCHAR(1),
    dropoff_longitude FLOAT,
    dropoff_latitude FLOAT,
    payment_type INT,
    fare_amount FLOAT,
    extra FLOAT,
    mta_tax FLOAT,
    tip_amount FLOAT,
    tolls_amount FLOAT,
    improvement_surcharge FLOAT,
    total_amount FLOAT
);

```

5 Data Load to Redshift

- Attached IAM Role: AmazonS3ReadOnlyAccess
- Loaded data using COPY:

```

COPY trips_even
FROM 's3://nyc-taxi-dataset-bucket/yellow_tripdata_2015-01.csv'
IAM_ROLE 'arn:aws:iam::535002864770:role/redshift_cluster_role'
CSV
IGNOREHEADER 1
TIMEFORMAT 'auto';

```

6 Query Workload Simulation

- Ran analytical aggregations and self-joins:

```

SELECT payment_type, COUNT(*), AVG(total_amount), SUM(tip_amount)
FROM trips_even
GROUP BY payment_type;

```

```

SELECT t1.payment_type, AVG(t1.total_amount + t2.total_amount)
FROM trips_even t1
JOIN trips_even t2
ON t1.VendorID = t2.VendorID
WHERE t1.passenger_count = t2.passenger_count
GROUP BY t1.payment_type;

```

7 Monitoring

- Monitored using CloudWatch and Query Editor v2:
 - Query execution time
 - CPU utilization by node
 - Slice usage via:

```
SELECT slice, COUNT(*)
FROM stv_blocklist
GROUP BY slice;
```

Observation

- ⚠ Uneven row distribution across slices
 - 🔴 One slice (node) had **disproportionate number of blocks**
 - ⚠ CloudWatch showed that one node had **spikes in CPU usage**
 - ✅ **DISTSTYLE EVEN** resulted in **data skew** under this dataset pattern
-

Next Phase: Optimization Using DISTSTYLE KEY (Planned)

Plan:

1. Create new table `trips_key` with `DISTSTYLE KEY` using high-cardinality column `VendorID`
 2. Reload same dataset into `trips_key`
 3. Rerun same query workloads
 4. Compare slice-level distribution and CPU balance across nodes
-

Conclusion (Pending Completion of Phase 2)

This simulation demonstrates that **DISTSTYLE EVEN is not always optimal**, especially when dealing with skewed keys. **Using a distribution key aligned with query filters or joins** is expected to improve both **execution time** and **node-level CPU utilization**.

Tools Used:

- Amazon Redshift (Provisioned Cluster)
 - Amazon S3
 - AWS CloudWatch
 - Redshift Query Editor v2
 - Kaggle NYC Taxi Dataset
-

Workspace Output (Paste Your Results Below)

 Paste screenshots, metrics, SQL outputs, and graphs here:

Screenshot of the AWS S3 console showing the 'nyc-taxi-dataset-bucket' page.

The left sidebar shows:

- Amazon S3
- General purpose buckets
 - Directory buckets
 - Table buckets
 - Access Grants
 - Access Points for general purpose buckets
 - Access Points for directory buckets
 - Object Lambda Access Points
 - Multi-Region Access Points
 - Batch Operations
 - IAM Access Analyzer for S3
- Block Public Access settings for this account
- Storage Lens
 - Dashboards
 - Storage Lens groups
 - AWS Organizations settings
- Feature spotlight [11]

The main content area displays the 'nyc-taxi-dataset-bucket' page with the following details:

- Bucket name: nyc-taxi-dataset-bucket
- Region: Asia Pacific (Mumbai)
- Last modified: June 17, 2025, 20:32:17 (UTC+05:30)
- Size: 1.8 GB
- Storage class: Standard

Object list:

Name	Type	Last modified	Size	Storage class
yellow_tripdata_2015-01.csv	csv	June 17, 2025, 20:32:17 (UTC+05:30)	1.8 GB	Standard

Screenshot of the AWS Redshift console showing the 'redshift-loadtest' cluster page.

The left sidebar shows:

- Amazon Redshift
- Redshift Serverless
- Provisioned clusters dashboard
- Clusters
 - Clusters
- Query editor
- Query editor v2 [2]
- Query and database monitoring [New]
- Datasources
- Integrations
 - Zero-ETL integrations
 - S3 event integrations [New]
- IAM Identity Center connections

The main content area displays the 'redshift-loadtest' cluster page with the following details:

General information:

Cluster identifier	Status	Node type	Endpoint
redshift-loadtest	Modifying	ra3.4xlarge	redshift-loadtest.cbnduwjtq5ur.ap-south-1.redshift.amazonaws.com:5439/dev
Custom domain name	Date created	Number of nodes	JDBC URL
-	June 18, 2025, 13:52 IST	4	jdbc:redshift://redshift-loadtest.cbnduwjtq5ur.ap-south-1.redshift.amazonaws.com:5439/dev
Cluster namespace ARN	Multi-AZ	Patch version	ODBC URL
arnaws:redshift:ap-south-1:535002864770:namespace:92c68612-f0ba-4caa-9238-6c87c1b2e11a	No	Patch 190 [2]	Driver=(Amazon Redshift (x64)); Server=redshift-loadtest.cbnduwjtq5ur.ap-south-1.redshift.amazonaws.com; Database=dev
Namespace register status	Cluster configuration	Storage used	
Deregistered	Production	0.00 of 488.3 TB used (0.0%)	

IAM > Roles > redshift_cluster_role

Identity and Access Management (IAM)

Search IAM

Dashboard

Access management

- User groups
- Users
- Roles**
- Policies
- Identity providers
- Account settings
- Root access management

Access reports

- Access Analyzer
 - Resource analysis
 - Unused access
 - Analyzer settings
- Credential report
- Organization activity
- Service control policies
- Resource control policies

CloudShell Feedback

Summary

Creation date: June 18, 2025, 13:50 (UTC+05:30)

Last activity: -

ARN copied: arn:aws:iam::535002864770:role/redshift_cluster_role

Maximum session duration: 1 hour

Permissions **Trust relationships** **Tags** **Last Accessed** **Revoke sessions**

Permissions policies (2) Info

You can attach up to 10 managed policies.

Filter by Type: All types

Policy name	Type	Attached entities
AmazonRedshiftFullAccess	AWS managed	1
AmazonS3ReadOnlyAccess	AWS managed	1

Permissions boundary (not set)

Generate policy based on CloudTrail events

© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

The screenshot shows the AWS IAM 'redshift_cluster_role' details page. The left sidebar contains navigation links for Identity and Access Management (IAM), Access management, and Access reports. The main content area displays the role's summary, including creation date (June 18, 2025, 13:50 UTC+05:30), last activity (none), ARN (arn:aws:iam::535002864770:role/redshift_cluster_role), and maximum session duration (1 hour). The 'Permissions' tab is selected, showing two attached AWS managed policies: 'AmazonRedshiftFullAccess' and 'AmazonS3ReadOnlyAccess'. There are buttons for 'Simulate', 'Remove', and 'Add permissions'. Below the policies, there is a section for generating a policy based on CloudTrail events. The bottom of the page includes standard AWS footer links for Privacy, Terms, and Cookie preferences.